

Jenny Ho

Culver City, 90232 | jennyho0221@gmail.com | (424) 660-4004 | www.linkedin.com/in/jenny-ho-32920a233

ML Engineer with research and production experience across LLM training, inference optimization, and on-device ML deployment. First author at NeurIPS 2024 and co-author at ACL 2025. Graduating from UCLA MEng in AI in August 2026; prior experience at Google and MediaTek Research. Seeking full-time ML Engineer roles with OPT authorization.

EDUCATION

University of California, Los Angeles (UCLA), Los Angeles, CA **Aug 2026**
Master of Engineering (MEng), Artificial Intelligence Track | GPA: 3.94/4.0
National Yang Ming Chiao Tung University, Hsinchu, Taiwan **Jan 2024**
Bachelor of Science in Computer Science | GPA: 4.19/4.30

SKILLS

- **Languages:** Python, C/C++, Java, JavaScript, SQL, Bash
- **ML/DS:** PyTorch, TensorFlow, TFLite, Hugging Face, scikit-learn
- **ML Systems:** LLM fine-tuning, training pipelines, inference optimization, evaluation/monitoring, experiment tracking
- **Infra:** Linux, Docker, Git, FastAPI, Node.js, parallel computing, Android ML deployment
- **AI/LLM:** Anthropic Claude API, LLM tool-calling, prompt engineering

PUBLICATIONS

- [1] "Stutter Makes Smarter: Learning Self-Improvement for Large Language Models," *Language Gamification Workshop 2024 @ NeurIPS, Dec 14, 2024, Vancouver, Canada*. (First Author)
- [2] "Let's Fuse Step by Step: A Generative Fusion Decoding Algorithm with LLMs for Multi-modal Text Recognition," *The Annual Meeting of the Association for Computational Linguistics (ACL)*, July 27 – Aug 1, 2025, Vienna, Austria. (Co-Author)

WORK EXPERIENCE

MediaTek Research, Taipei, Taiwan **May 2024 — May 2025**
AI and Machine Learning Engineer, AI Research Center

- Led research and implementation of "Stutter Makes Smarter," a method to improve small LLMs (160M–1B parameters) by storing and reusing hidden states across passes; introduced new attention parameters to stabilize training; accepted at NeurIPS 2024 Workshop. (First Author)
- Re-architected training pipeline with parallel data preprocessing and async prefetching, overlapping I/O with compute to reduce GPU idle time and overall compute cost by ~20%.
- Collaborated on the Generative Fusion Decoding framework, refactoring and open-sourcing the codebase to improve reproducibility; reran evaluations across benchmark datasets to validate consistent performance; work accepted at ACL 2025.
- Refactored and modularized a Hugging Face-based ML codebase into reusable libraries and deployment-ready packages, streamlining model packaging, serving, and reproducibility across teams.
- Developed evaluation and monitoring dashboards for model accuracy drift and inference latency; integrated periodic checks into deployment workflows to prevent regressions.

Google, New Taipei City, Taiwan **June 2023 — Sept 2023**
Software Engineer Intern, Pixel Recorder Team

- Built and extended internal QA tooling for the Pixel Recorder team, adding new test features to the Validation tool covering audio post-processing pipelines; adopted by both QA engineers and internal developers.
- Automated batch audio testing with Python scripts, replacing a manual QA workflow and reducing repetitive test execution time across the team.
- Redesigned the Validation tool UI to support audio file selection, real-time progress tracking, and metrics visualization, improving test result visibility for QA workflows.

Gamania Digital Entertainment, Taipei, Taiwan **July 2022 — Aug 2022**
AI Research Intern, Big Data Research Center

- Built recommendation pipeline and feature engineering workflows supporting personalization; improved CTR by 15% through iterative hyperparameter tuning validated with offline experiments.

PROJECTS

AI Financial Assistant **Mar 2026 — Present**

- Designed and built a full-stack AI financial assistant with a Claude-powered agentic tool-calling loop, routing user queries across six tools covering real-time stock data, financial statements, risk metrics, and macroeconomic indicators via yfinance and FRED API.
- Architected the backend in FastAPI with modular tool definitions and a multi-turn agentic loop; built a Next.js frontend with a chat interface and a portfolio performance chart visualizing normalized historical returns.

Reinforcement Learning for Board Games **Feb 2024 — Mar 2024**

- Implemented tabular Q-learning and DQN agents for a turn-based board game with ϵ -greedy exploration, self-play/heuristic opponents, and legal-move masking.
- Built a training pipeline with experience replay (minibatch sampling) and target network updates; ran ablations on ϵ schedules and replay size, tracking win-rate and learning curves.